

Genetic Mutation Classification based on Clinical
Evidence to Enable Personalized Medicine for Cancer
Treatment

By

LIU Kai Yang, Louis
(1830004016)

A Final Year Project thesis (STAT4004)
submitted in partial fulfillment of the requirements
for the degree of

Bachelor of Science (Honours)
in Statistics

at

BNU-HKBU
UNITED INTERNATIONAL COLLEGE

December, 2021

DECLARATION

I hereby declare that all the work done in this Project is of my independent effort. I also certify that I have never submitted the idea and product of this Project for academic or employment credits.

LIU Kai Yang, Louis
(1830004016)

Date: _____

Genetic Mutation Classification based on Clinical Evidence to Enable Personalized Medicine for Cancer Treatment

LIU Kai Yang, Louis

Science and Technology Division

Abstract

Traditionally, interpretation of the clinical documents to diagnosis the genetic variants type for patients with cancer is very time-consuming, it will cost lots of human efforts. In recent years, the emergence of Machine Learning and Deep Learning classification models that combined with Natural Language processing techniques greatly facilitate the diagnosis process based on the clinical texts. In the future, personalized medicine for cancer treatment may become true by using advanced gene classification models. In this project, our goals is to classify 9 genetic mutation classes based on the clinical evidence. There are two proposed models. For the Word2Vec + LightGBM models, its computational strength with the balance of accuracy help it reach a 60% of accuracy and got a desirable 2.52648 private score in Kaggle. For the Word Embedding + Bidirectional GRU with Attention model, due to the strength of gate structure in GRU, full picture of text information retrieved by bidirectional layers, and also the contribution of attention mechanism in focusing more important part of texts, we got a great improvement and get an 85% accuracy result and finally reach 2.36962 private score with rank 88 out of 1386 in Kaggle.

Keywords: LightGBM, GRU, Bidirectional, Attention, Word2Vec

Preface

A Genetic Mutation Classification Project based on Clinical evidence. Processed Natural Language Processing and use Machine Learning methods LightGBM as well as Deep Learning models Bidirectional GRU with Attention.

Acknowledgment

This project would not have been possible without the support of many people. Thanks to my advisor Dr. Yuhui Deng for holding so many seminars, providing great support, and giving instructions during my final year project. Also, I will give many thanks to my friends in the same group who offer guidance and supports.

Contents

1	Introduction	1
1.1	Background information	1
1.2	Motivation and Project Objectives	2
1.3	Related Work	3
1.3.1	Machine Learning in Classification	3
1.3.2	Deep Learning in Classification	5
1.3.3	Discussion	7
1.4	Models Overview Diagram	8
2	Data Exploration and Analysis	9
2.1	Dataset Descriptions	9
2.2	Data Analysis	10
2.3	Text Preprocessing	13
3	Methodology	16
3.1	Overview	16
3.2	Word Embedding	17
3.2.1	Word2Vec + Average Feature	17
3.2.2	Word Embedding from Keras	19

3.3	Model 1 - LightGBM	20
3.3.1	Background information	20
3.3.2	Basic concept and Algorithm	22
3.4	Model 2 - Bi-GRU + Attention	26
3.4.1	Basic Concepts in Neural network	26
3.4.2	Recurrent Neural Network	30
3.4.3	Gate Recurrent Unit	31
3.4.4	Bidirectional Layers	33
3.4.5	Attention context	33
3.4.6	Final Proposed Model	35
4	Experiment and Results	37
4.1	Word2Vec + LightGBM	37
4.2	Word Embedding + Bi-GRU,Attention	40
4.3	Results	44
5	Conclusions	47
A	Python Code	49
	Bibliography	62

List of Tables

4.1	Parameters setting for Word2Vec and Average Features . .	38
4.2	Parameters setting for LightGBM	38
4.3	Training for LightGBM	39
4.4	Model Evaluation for LightGBM	40
4.5	Parameters for Bi-GRU with Attention Model	41
4.6	Bi-GRU with Attention Model	42
4.7	Training for Bi-GRU with Attention Model	43

List of Figures

1.1	Project Overview Diagram	8
2.1	Distribution of Genetic Mutation Classes	11
2.2	Top 10 genes with maximal and minimal occurrences . . .	11
2.3	The distribution of gene in different classes	12
2.4	Word Cloud	13
2.5	Statistics description for Word count	14
2.6	Text length distribution	15
3.1	CBOW and Skip-Gram Structure	18
3.2	Word Embedding	19
3.3	A Simple Decision Tree Model	21
3.4	Gradient Boosting Decision Tree Model	22
3.5	Histogram-Based Algorithm	23
3.6	A Simple neural network	26
3.7	A Simple MLP network	28
3.8	A Simple RNN structure	30
3.9	The structure of GRU	32
3.10	RNN with bidirectional layer	34
3.11	The structure of Attention Context	35

3.12	Structure of Bi-GRU+Attention Model	36
4.1	Classification Results-1	44
4.2	Classification Results-2	45
4.3	Classification Results-3	45
4.4	Final Results	46

Chapter 1

Introduction

1.1 Background information

According to International Agency for Research on Cancer, nearly ten million death in 2020 make cancer continually become the disease that dominates death around the world.[1] Cancer is caused by genetic mutation and normal cells thus transform into tumor cells. Since tumors always contain many types of genetic mutation and due to the mutational heterogeneity, different subgroups have different sensitivity to chemotherapy drugs. Although the use of chemotherapy drugs may temporarily suppress tumors, the selective pressure formed by them makes the drug-sensitive subgroups gradually disappear and become insensitive. The subgroup reproduces and causes tumor recurrence or metastasis. Therefore, more attention should be paid to heterogeneity to enable personalized cancer treatment.[2] That is also to say, identification and classification of the particular type of gene mutation from the clinical documents that cause cancer are essential. However, this process suffers from multiple drawbacks.[3] First, the

identification of genetic mutation requires professional knowledge in gene and medicine science, and these constraints shut many ordinary labors out. Also, the interpretation of those clinical texts is based on the subjective perspective from person to person even though they are professional, and it always depends on the situation due to the huge difference between patients. Moreover, plenty of works are required for genetic mutation classification if the doctor is dealing with daunting denotation work manually since it is complicated and extremely time-consuming.

1.2 Motivation and Project Objectives

Tumor diagnosis and treatment always require precisely clinical documents. Also, distinguishing between mutations that promote tumor growth (drivers) and neutral genetic mutations (passengers) remains a challenge.[4] As a result, most interpretations of clinical literature still need to be handled manually.[5] Since this process is less efficient and subjective, new techniques, thanks to the development of NLP (Natural Language Processing) techniques, have emerged to address the gene classification problems among clinical evidence.[6] Generally, clinical document research mainly focuses on the words used (medical concepts, notations), sentence connections, and semantic attributes (describe the condition of patients). Although some of the studies try to put more effort to evaluate patients or population level, the evaluation approaches are often inconsistent due to the difference of objectives and methodology priorities.[7] Therefore, to face all the challenges mentioned before, applying different NLP techniques to the clinical document, and finding out some appropriate methodologies

to automatically classify the genetic variations become critical and valuable. In the future, with the development of the application of machine learning and deep learning methods in clinical text, personalized medicine will be facilitated and benefit cancer treatment.

The ultimate purpose of this project is to develop some classification models to give the classified result of the gene mutation classes and realized personalized medicine for cancer treatment. Based on the clinical evidence, different genetic mutation types are asked to be classified by referring to the information given in the gene variation document. To achieve this goal, a powerful with high accuracy machine learning or deep learning model is required.

1.3 Related Work

Many studies have done in-depth analysis on the classification of genetic mutation texts in medical texts. Nowadays, the most popular classification model can be divided into machine learning method and deep learning method. Since the twentieth century, many new models are developed so that they can handle the classification of natural language better.

1.3.1 Machine Learning in Classification

With the development of the computation, instead of dealing with the clinical documents manually, the emergence of the traditional machine learning method can be seen as a great improvement.[8] Mark Singh et al. used a naive Bayesian classifier, which was widely used in the early 1960s, to accurately screen clinical reports and got detect abnormal radiology

results. They conclude that it can help doctors review radiology reports more effectively.[9] In the 1990s, a powerful model SVM was developed quickly and is still popular as a classification model by now. Adam Wright et al. in 2013 successfully get a good result by using a Support Vector Machine-based classifier by identifying EHR (Electronic Health Record) progress documents that are related to diabetes.[10]. In 1995, tree structure based classification methods came out, where random decision forests proposed by Tin Kam Ho increase the generalization accuracy for tree-based classifiers.[11]. Then, followed by the development of the Gradient Boosting Decision tree that was proposed by Jerome H. Friedman from Stanford university in 2001, the tree-based method greatly advanced.[12]. However, the cons for GBDT is that the computation cost is too high when facing modern industry data. Hopefully, an optimized GBDT based method XGboost was proposed in 2016, it greatly increases the speed and performance due to its improvement on loss function, avoiding overfitting, and computing.[13] Gupta et al. did a great work by using several machine learning models like random forest and XGboost to do the gene classification and reach a good results.[5]

However, in some situations, the traditional machine learning methods also get low accuracy. More advanced machine learning techniques have been proposed to face those challenges. The emergence of LightGBM (Light Gradient Boosting Machine) in 2017 proposed by Microsoft has great improvements in dealing with clinical documents.[14] It is intended to design for satisfying the needs of the industry to pursue a higher accuracy and shorten the calculation time of the model. It can also ensemble multiple learning classifiers to deal with more complex literature. Xuan Qin et al. proposed an ensemble Light GBM learning method, which

integrates multiple BERT (Bidirectional Encoder Representations from Transformers) models to improve the classification efficiency of titles and abstracts filter.[15] The drawbacks of machine learning techniques is also obvious. Feature extraction done by experts in the medical field is needed and thus it will take plenty of labor cost.[16]

1.3.2 Deep Learning in Classification

Deep learning can be seen as an essential branch of machine learning. Instead of extracting features from the data set, which is part of the procedure in machine learning, deep learning can automatically learn those features and classify them to get appropriate results and save a huge amount of time.

Based on the original multiple layer perceptron, RNN (Recurrent Neural Network) was proposed in the 1980s that intend to solve the data with time and sequence. Meenu Gupta et al. also applied the Word2Vec text transformation model and use RNN (Recurrent Neural Network) model and finally get a well-performed result than other proposed classifiers.[5] Later in 1997, a modified type of RNN, LSTM (Long Short-Term Memory) was designed for combining the long and short-term memory by using the gate structure, and also solving the vanishing gradient problems.[17] Tang et al implement a modified Bi-LSTM-CRF model to do de-identification of the clinical documents and reach a great result. In 2014, GRU (Gate Recurrent Unit), a new type of RNN model that makes some modifications based on LSTM, was proposed to improve the computation performance and keep a relatively good result. In the same year, Attention context techniques were used in RNN in dealing with the images to make attention

popular.[18] It enables the neural network to focus more on the important information while lowering the weight with the meaningless part. Later, many recurrent models combined with attention were proposed in solving the text classification.

For the development of CNN, though it is originally applied in images, some research based on medical documents is also finished. When Yujia Bao et al. study the Medical documents that are related to cancer susceptibility genes, CNN (Convolutional Neural Network), as one of the important deep learning methods are implemented. They conclude that it can help doctors and researchers get to acquire the most recent updated knowledge among medical literature in the field of gene mutation.[19]. Also, in Yanshan Wang et al. works, CNN (Convolutional Neural Network) express more advantages in proposing a clinical text classification paradigm. It shows a higher accuracy rate compared to Random Forest, SVM (Support Vector Machine), and MLPNN (Multilayer Perceptron Neural Networks).[20]

In 2018, a new deep learning method called BERT (Bidirectional Encoder Representations from Transformers) was proposed by Google. The unique feature of BERT is its deep bidirectional architecture, which only requires an additional output layer to fine-tune the pre-training BERT to meet various tasks, and there is no need to modify the model for specific tasks. [21] Later, in Yuhan Su et al. research, they applied a modified BERT model on genetic mutation classification and come up with the conclusion that BERT is applicable in dealing with clinical evidence and can contribute a lot in speeding up the diagnosis and treatment of cancer. What is more, the more precise and personalized medicine treatment will be facilitated.[22] However, in Stephen Wu et al. methodical review of the

application of deep learning techniques among clinical evidence, they point out that there is still a long way for deep learning to go in the future.[23]

1.3.3 Discussion

From the related work above, we can see that there are some models with good performance in both machine learning and deep learning method. For example, LightGBM shows a high performance when dealing with the clinical texts, the RNN model is good at solving the text with sequential information, BERT achieves a high accuracy when classifying the gene variants. All of the models have its strength and weakness, therefore, how to appropriately use different techniques to achieve better results become critical.

1.4 Models Overview Diagram

In this project, two models will be proposed to solve the gene variants classification problems. For the first machine learning model, we will use word2Vec + LightGBM. For the second deep learning model, we will use word embedding + Bidirectional GRU with attention context. The whole process of this project is listed below in Figure 1.1.

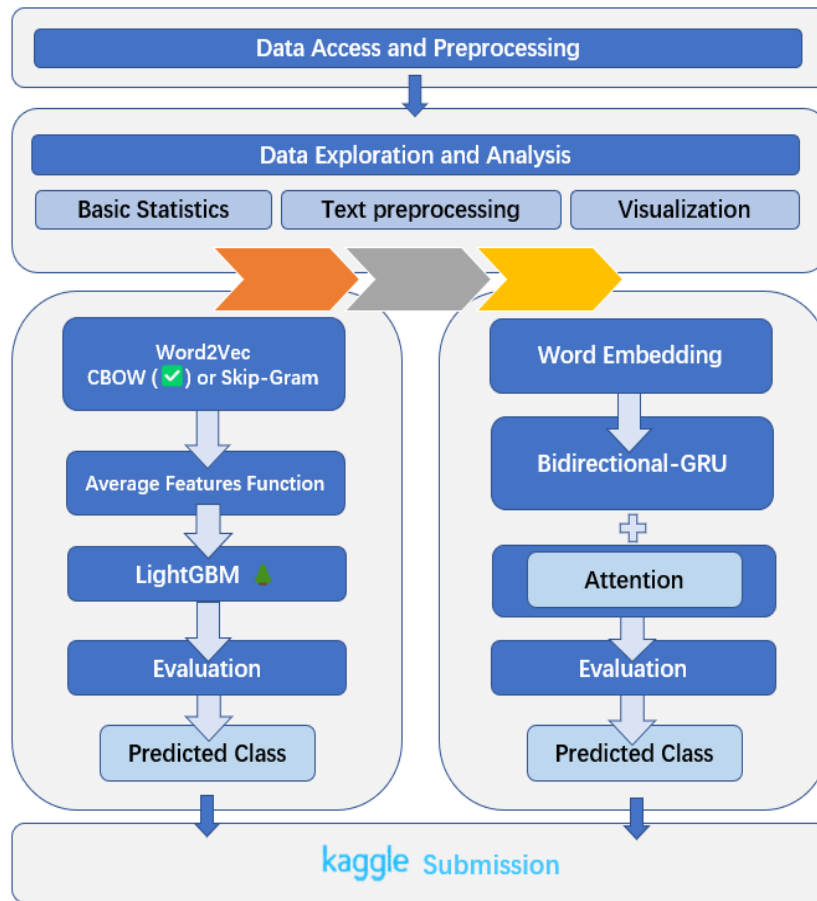


Figure 1.1: Project Overview Diagram

Chapter 2

Data Exploration and Analysis

2.1 Dataset Descriptions

In 2017, Memorial Sloan Kettering Cancer Center (MSKCC), which is the most renowned hospital in the tumor treatment area, released a competition that focus on genetic mutation classification to achieve personalized medicine treatment. The dataset they provide is all manually annotated by the doctor with proficient knowledge in cancer treatment. There are two different kinds of documents, clinical text and variation type of different genes. Also, according to the different combinations of genes and variations, the annotated classes are given. The texts are the clinical records or evidence are written by the doctor when diagnosing their patients and the gene variations are marked by those doctors based on their expertise skills.

There are two types of CSV files in this project. One of them contains clinical texts and the other one contains the genes and variation information.

For the training set, the file that contains the genetic mutation classification has four features in total. "ID" work as the identity of information for each corresponding text in another file, "Gene" gives the information of which specific genetic mutation is located, "Variation" provide the different type of mutation, and "Class" are the label of the final classification of this genes mutation. Also, Nine classes in total are provided, and our goal is to predict the class result by using our models. For the file containing the clinical texts, there are two features in total. The first one is "ID", which is the same as the first file. Another feature is the clinical evidence for each ID, which contains 3321 observations in total. For us, we need to do natural language processing to deal with those texts and train our model to get the final classification results.

For the testing dataset, there are 5668 observations in total and all the information is the same as the training dataset except for the unlabeled classes. This dataset will be used as the validation set in this project.

For the Stage 2 testing dataset, though all the information is the same as the testing set, which does not contain the classification labels, there are only 986 observations, and this data set will be used to predict the final prediction results. And we can submit the final classification result on Kaggle to check the score and loss.

2.2 Data Analysis

Before we go through our text data, let us do some data exploration so that we can have an overview of our dataset. Since there are Nine classes of genes variation, we would like to see the frequency of each class by plotting the distribution of genetic mutation classes in a histogram. From




Figure 2.1: Distribution of Genetic Mutation Classes

Figure 2.1, we can see that Nine classes are shown, Class 7 have the highest frequency and class 3, 8, 9 have a relatively lower frequency.

Genes with maximal occurences		Genes with minimal occurences	
Gene		Gene	
BRCA1	264	KLF4	1
TP53	163	FGF19	1
EGFR	141	FANCC	1
PTEN	126	FAM58A	1
BRCA2	125	PAK1	1
KIT	99	ERRFI1	1
BRAF	93	PAX8	1
ALK	69	PIK3R3	1
ERBB2	69	PMS1	1
PDGFRA	60	PPM1D	1

Figure 2.2: Top 10 genes with maximal and minimal occurrences

There are not only nine kinds of gene mutation classes, different genes and mutation types combine to decide which specific 1 out of 9 classes there are belong to. After statistics, the number of unique genes is 1522,

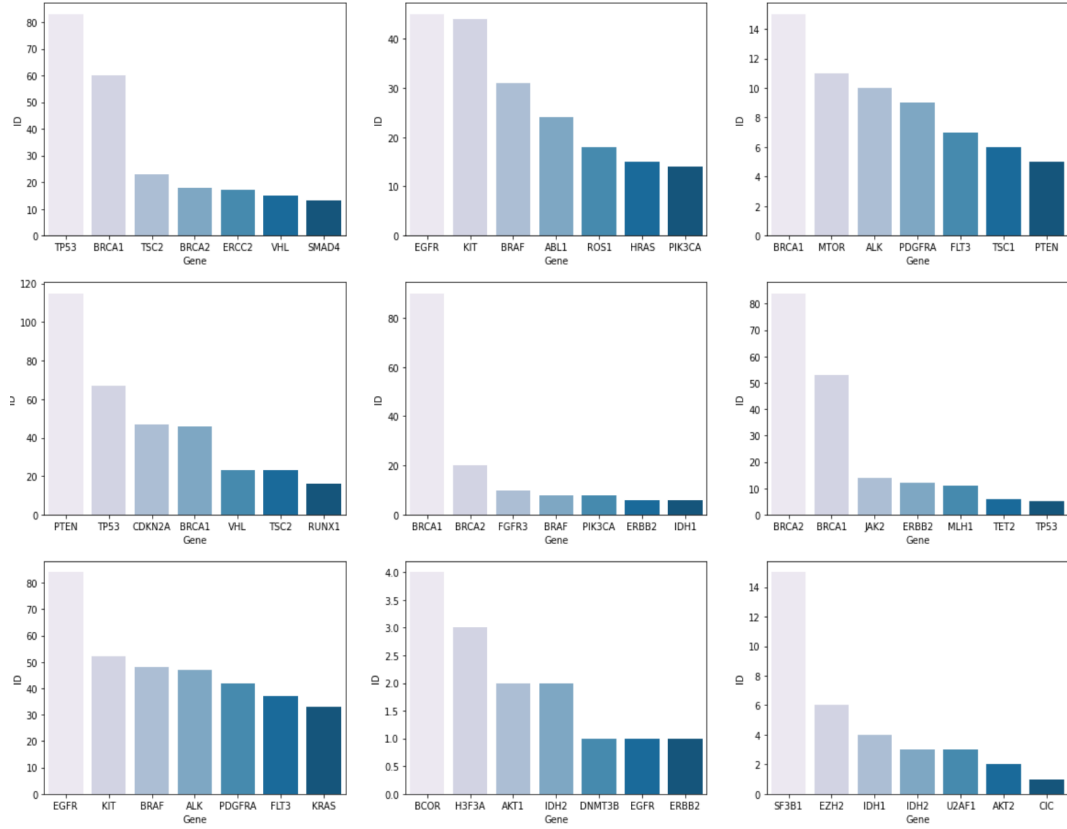



Figure 2.3: The distribution of gene in different classes

the number of unique variations is 347. Figure 2.2 shows the top 10 genes with maximal occurrences as well as minimal occurrences so that we can have a view of the number of the type of the most significant gene. And figure 2.3 shows the distribution of genes in 9 different classes. We can get the conclusion that some of the genes are highly dominating in their class. For example, Gene BRCA1 has the highest frequency and dominates in class 5.

When dealing with text data, text preprocessing is always the most important part that we should consider. Firstly, tokenization is applied to separate the sentence from words. Then, the removal of punctuation and lower casting have proceeded. Also, some relatively meaningless stop words are removed. For example, "a", "and", "but" and so on.



Some word statistics are also implemented, including word count, text length, and word-cloud plot to show the frequency of the words. From figure 2.4, we can find that words like "patient", "cell", "line", "amino", "acid" are frequently used in our clinical documents. Besides, text length distribution is also visualized to see the text count for each class. Also, figure 2.5 shows the statistics description for Word count. From figure 2.6, we can see that some texts in class 7 have a huge amount of words compared to the classes like 3,4,5,6, and 8. The average word count for class 7 is about 11000 words, and the highest one reaches 76733-word counts. Besides, most of the text in class 3 is around 7000 words are also observed.

	count	mean	std	min	25%	50%	75%	max
Class								
1	568.0	9441.841549	6511.773899	1.0	4969.75	7302.0	12866.25	52918.0
2	452.0	9304.159292	7621.158837	116.0	4185.00	6808.0	12219.50	61945.0
3	89.0	6749.213483	3712.931889	1737.0	4283.00	5571.0	7409.00	27290.0
4	686.0	8975.769679	7270.444322	53.0	4560.00	6351.0	11536.25	43812.0
5	242.0	7504.384298	3895.755024	183.0	5245.00	6426.0	9513.00	24130.0
6	275.0	7177.952727	3833.400979	1.0	4498.50	6587.0	7847.00	24519.0
7	953.0	11433.295908	10104.998688	1.0	4871.00	8254.0	14592.00	76733.0
8	19.0	10809.368421	5645.232888	2111.0	5586.00	11248.0	15529.00	20615.0
9	37.0	12795.675676	10208.050296	1147.0	4937.00	10910.0	15791.00	45078.0

Figure 2.5: Statistics description for Word count




Figure 2.6: Text length distribution

Chapter 3

Methodology

3.1 Overview

There is two classification model that will be proposed in this project. The first one is the machine learning method LightGBM (Light Gradient Boosting Machine), which is an improved model of GBDT (Gradient Boosting Decision Tree model). The second one is the deep learning method RNN (Recurrent neural network). We will choose GRU (Gate Recurrent Unit) cells with a bidirectional layer and add attention context techniques as our final model. In this section, we will discuss these two models from the very beginning in a more understandable way, and show how these two models are applied to our classification.

3.2 Word Embedding

Since we cannot simply use the natural language as the input of the model, some transformation is needed so that we can let the computer understand the data that we are going to use before the classification.

In this project, for Model 1 - LightGBM, the Word2Vec Model is implemented and we use the average feature vector function to get the matrix representation of the words from word2vec and use it as the input of LightGBM. For Model 2 - Bi-GRU + Attention Model, we use the word embedding model from Tensorflow, where the input of model 2 is a tensor with multiple layers of the matrix representation of words.

3.2.1 Word2Vec + Average Feature

Word2Vec is a prevalent model with many applications. For example, it can be applied in a classification problem, clustering problem, calculation of word similarity, and so on. In this section, we will briefly introduce how the CBOW (Continuous Bag of Words) model and Skip-gram (Continuous Skip-Gram) model works in generating the word vector that will be used in the classification model.

CBOW and Skip-gram have different ideas when predicting the word vector. For a CBOW model, the input is a continuous bag of words from a sentence except for the word that we want to predict. We use the vector form of the words as the input and get the predicted output (also in vector form) through a neural network. Similarly, for a skip-gram model, the input is a central word after encoding, and through the neural network to get the predicted words in the sentence except for the central word. For the structure of the neural network, we will explain it further in the

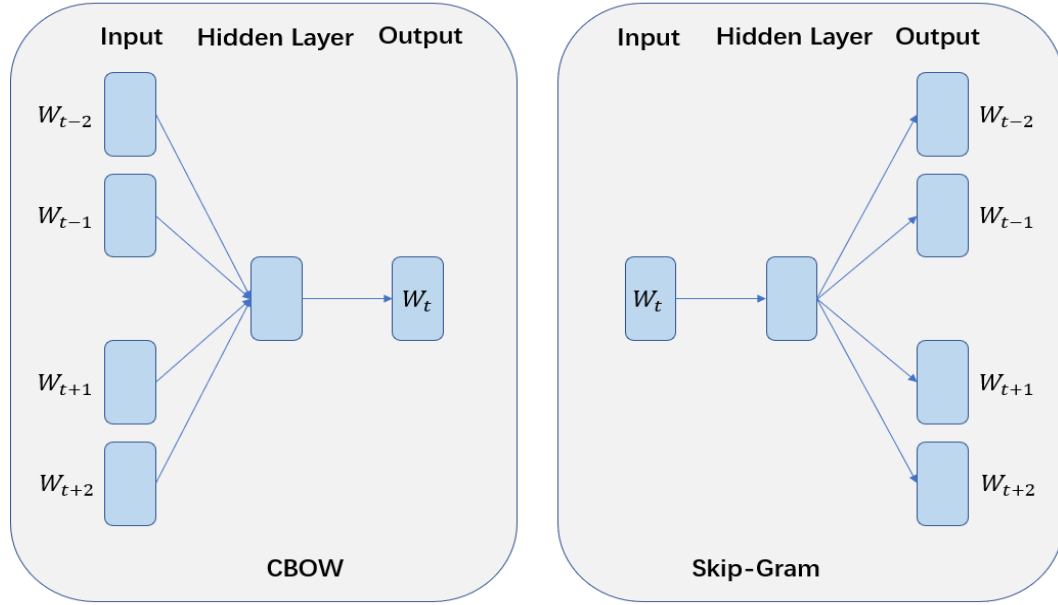


Figure 3.1: CBOW and Skip-Gram Structure

following chapter when we talk about the multiple layer perceptron.

In a classification problem, the input is our texts, which contain many sentences and paragraphs, we use those texts as the input of the Word2vec model and get the word vector from the model parameter is the vector representation of the texts. Here, we need to know that the word vector we get from this Word2Vec model is the weight matrix of the network. When we get the vector representation of the text, we find that if we would like to use the LightGBM as the model, we need to make sure that the input size of the word-represented matrix should equal to (text length x embedding dimension) so that it fits the classification label in the training set. Therefore the average feature function is used so that we can generate an input matrix with the correct size by lowering the number of

word vectors in the weight matrix to the length of the text. For example, we set the word embedding dimension, which is the dimension of the hidden layer to be 200, and the shape of the weight matrix we get from the input side is 5852 times 200. Since there are 3321 rows of text data in our training set and there is 3321 corresponding label that we need to train. So What the average feature function do is to shorten the 5852 rows in the weight matrix to 3321 rows by calculating the average of the word vector to make sure the input size of the LightGBM model is correct.

3.2.2 Word Embedding from Keras




Figure 3.2: Word Embedding

The representation matrix after the average feature function has some weaknesses since it ignores so much important information from the text. When we are using a deep learning method that we propose in model 2, we generate a vector matrix with multiple layers to keep more information from the text.

To explain it further, there are two steps to get the result. First,

we should use the Tokenizer to change texts to sequences, which means for each training sample, we will get a corresponding sequence. Since each texts sample has different words, the sequence size for each of them may not be equal. Therefore, we need to do the second steps, which is padding. Based on the sequence length that we would like to set, we will pad those sequences to be the same length. Then, based on the total vocabulary size, pre-set embedding dimension, and the sequence length for each training sample, we can finish the word embedding. Notice that the pre-set embedding dimension is the dimension that we would like to give for each word in the sequence. Therefore, the size of the vector representation of a text should become cubic in three dimensions, where the shape should be the number of observations in the training sample times the sequence length times the embedding dimension. Figure 3.2 shows the process more clearly.

3.3 Model 1 - LightGBM

3.3.1 Background information

LightGBM is a kind of GBDT (Gradient Boosting Decision Tree) model. Therefore, to have a better understanding of LightGBM, we can briefly introduce the concept of Gradient Boosting Decision Tree first.

To begin with, we need to know that GBDT belongs to the decision tree model. For a classification problem, the decision tree, which has a top to a button tree structure, has many criteria in their nodes, and those criteria in the parents' node can split the data into left child node and




Figure 3.3: A Simple Decision Tree Model

right child node to achieve the classification purpose. While the gradient boosting decision tree is an Ensemble Learning method in machine learning. Generally speaking, for a classification problem, we need to calculate the log-likelihood for each observation as the first step, generating tree models that fit the data and calculating the residuals for each of the observations. Then, update the predicted result and reiterate the above procedures. The basic structure of the model is shown in figure 3.4. We will write down the GBDT algorithm for multiple classification problems to explain it further.

After talking about the GBDT model, we can go further to introduce the LightGBM model. A common machine learning method, or deep learning method, has less restriction when training the huge amount of data from the industry by using the mini-batch method. However, for a GBDT model, since it applies a level-wise tree growth algorithm, it is required to go through the whole data set for each iteration, which will




Figure 3.4: Gradient Boosting Decision Tree Model

lead to resource exhausted problems for the computer. Also, it is time-consuming to split the data for several training steps in the industrial application. Therefore, a highly efficient model LightGBM that applies a leaf-wise tree growth algorithm came out aiming to solve those problems.

3.3.2 Basic concept and Algorithm

In this section, we will go through the main idea of the Four important algorithm proposed in the LightGBM to see why it is a good model compared to the GBDT. However, since our project is focused on the application of LightGBM as a real classification problem, we will talk more about the understanding of the algorithm instead of how the coding works inside LightGBM.

Histogram-based Algorithm

Generally speaking, the design of this histogram-based algorithm is to make the storage of data easier and let the computation more efficient.

Also, the model will become more robust. The first step of the histogram algorithm is to determine how many bins are required for each feature and assign an integer number to each box. Then, we need to divide the range of floating points into several intervals, where the number of intervals should equal the number of the bin. Third, update the sample value in the bin to the new value corresponding to each bin. Finally, we have done the simplified of large-scale data to a histogram-based dataset. At this time, the features have become discrete and easy to store and have a low cost of computation. For the weakness of this algorithm, the splitting point of the decision tree is not as easy to find as GBDT. However, in the structure of gradient boosting, the precise of the splitting point is not so important, even in this case we can avoid over-fitting to some degree.

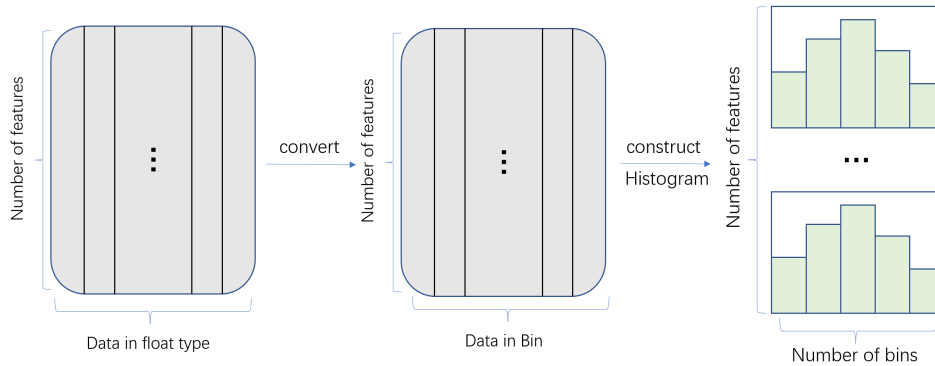


Figure 3.5: Histogram-Based Algorithm

Gradient-based One-Side Sampling

The motivation for using the Gradient-based One-Side Sampling is to have a compromise between the learning accuracy for the decision trees and reducing the size of the data set.[14] To be more specific, we would like to reduce the samples with the smaller gradients and use the remaining part of the samples to calculate the information gain in the LightGBM. Here is the problem we may face, samples with small gradients have relatively small training errors, indicating that the data has been learned by the model very well. However, if we discard this part of the data with small gradients directly, we will change the distribution of the data and will affect the accuracy of the training model. To avoid this problem, the GOSS algorithm is proposed.

When we are doing sampling, it is expected to discard those data that contribute less to the information gain. Since samples with large gradients have more impacts on the information gain, the GOSS Algorithm first sorts all the values of the features that are to be split in descending order of the absolute values. Then, select the top $a \cdot 100\%$ samples. After that, $b \cdot 100\%$ samples with a lower gradient will be randomly selected and will be amplified by multiplying with a constant $(1-a)/b$ in the calculation of information gain. As a result, the distribution of the original sample data with a lower gradient will be discarded with less cost and the learning accuracy will not be affected so much at the same time.

Greedy Bundling

Greedy bundling aims to reduce the number of features to increase computation efficiency. The main idea of greedy bundling is to shrink the dimension of the features by bundling them together. In the high-dimensional data, especially when we are doing the encoding, sparse matrix is a very common form. However, we know the reason that we call it a sparse matrix is that it has great size but with less information. It is easy to find some of the features that are mutually exclusive so that we can bundle them together to reduce the dimension and optimize the time cost for the algorithm. The greedy bundling algorithm solves the exact problem about which kind of features can be bundled.

Merge Exclusive features

For merge exclusive features algorithm, it completes the remaining jobs of the greedy bundling algorithm. When we figure out what features should be bundled together, we need to construct the bundling part by using a new algorithm. What needs to be mentioned after merging is that we need to make sure the data in the bundled features can still be recognized. For the first algorithm, since we have already transformed the float data into different bins, it is a good idea for us to make the merged features set in the bins. To achieve this goal, we can add some offset on the original data to ensure that the range of two exclusive features is merged safer.

3.4 Model 2 - Bi-GRU + Attention

3.4.1 Basic Concepts in Neural network

In this section, we will talk about some basic concepts in a neural network so that we can have the prerequisite knowledge to understand Recurrent Neural networks and their improved version.

Perceptrons and Sigmoid




Figure 3.6: A Simple neural network

To begin with, perceptrons can be seen as the fundamental of a neural network. The basic structure of perceptrons can be described as a transformer from the input to the output. Let's say a linear model, $X_1, X_2, X_3, \dots, X_n$ is the input of the perceptron, b is the bias, and Z is the output. What needs to be mentioned is that we can give the input with different weights to decide which of them are more important that contribute to the result Z . Therefore, we can write a simple linear equation 3.1. We can go further by setting a threshold. If the output Z is larger

than a threshold that we set, we label them with 1, otherwise, label it with 0. At this time, we finish a simple classification model.

$$Z = \sum_{i=0}^n W_i \cdot X_i + b \quad (3.1)$$

For the sigmoid, it makes some small changes on the output Z from equation 3.1 by using the sigmoid function 3.2. The sigmoid function is continuous, monotonic increasing, and differentiable. It enables us to get a range from 0 to 1 to represent the probability so that we can make some connection to some distribution. Also, a great propriety show in equation 3.3 that the derivative of a sigmoid function is equal to itself multiplied by one minus itself, which can help us calculate its derivative more easily. In the neural network, we always need an activation function for each neuron, which can introduce some kind of nonlinear factor to face the data in the real world. And the sigmoid function is one of the most common use ones.

$$a = \sigma(Z) = \frac{1}{(1 + e^{-Z})} \quad (3.2)$$

$$\sigma'(Z) = \sigma(Z) \cdot (1 - \sigma(Z)) \quad (3.3)$$

The basic structure for Multiple Layer Perceptrons

Let's move further to see the basic structure for a multiple-layer perceptron neural network as shown in figure 3.7. We will explain the structure and introduce the forward-propagation as well as the main idea of back-propagation. We can see that this network contains 3 layers in total, the

input layer, the hidden layer l , and the output layer L . In reality, there may be many layers l in the hidden layer, but we will use only one for explanation purposes.




Figure 3.7: A Simple MLP network

For forward propagation, our aim is to get the output of the network and calculate the cost. The most left part $X_1^1, X_2^1, \dots, X_j^1$ are the input layer of this network. W_{kj} are the weights from the first input layer to the next layer. Different from the last section, this time we write it in matrix form for simplicity and noticed that the shape for a weight matrix should contain the number of the neuron in hidden layers in a row and the number of the neuron in input in columns. The hidden layer in a neural network is always very flexible since we can change the number of layers and their dimension. But again we need to meet the dimension

requirement for matrix multiplication. In the hidden layer l in figure 3.7, we can see that $Z_1^l, Z_2^l, \dots, Z_k^l$ are the input in this layer, and Z^l in vector form should equal to $W_{kj}^l \cdot X^1 + b^l$, where b_l is the error in the layer. After activation we can get the output $a_1^l, a_2^l, \dots, a_k^l$. Then, the same procedure as the last steps, we multiply the output a from the hidden layer and the weighted matrix W_{nk} to get the input for the output layer L $Z_1^L, Z_2^L, \dots, Z_n^L$, and their vector form $Z^L = W_{nk}^L \cdot a^l + b_L$. After activation, we can get $a_1^L, a_2^L, \dots, a_n^L$. In the final step, we can calculate the output C by using the cost function defined by $\frac{1}{2} \|y - a^L\|$, where y is the real labeled output.

For the backward propagation, the main idea is to figure out how the changing of the weights and biases in our network can change the cost function, which is $\frac{\partial C}{\partial W^l}$ and $\frac{\partial C}{\partial b^l}$. Based on that, we can update the Weights and biases by multiplying the learning rate to achieve better results. To calculate the partial derivative in a convenient way, we define an intermediate error $\delta_j^L = \frac{\partial C}{\partial Z_j^L}$, which is the output error from the layer l . Based on the chain rule step by step, we have the four fundamental equations for backward propagation.

$$\delta^L = \nabla_a C \odot \sigma'(Z^L) \quad (3.4)$$

$$\delta^l = ((W^{l+1})^T \delta^{l+1}) \odot \sigma'(Z^l) \quad (3.5)$$

$$\frac{\partial C}{\partial W_{jk}^l} = a_k^{l-1} \delta_j^l \quad (3.6)$$

$$\frac{\partial C}{\partial b^l} = \delta_j^l \quad (3.7)$$

In this section, we will not discuss the proof of those equations, but we need to get familiar with those equations so that we know how the rate of change of weights and biases with respect to cost works, which will provide us with the fundamental knowledge to understand the Recurrent Neural Network.

3.4.2 Recurrent Neural Network

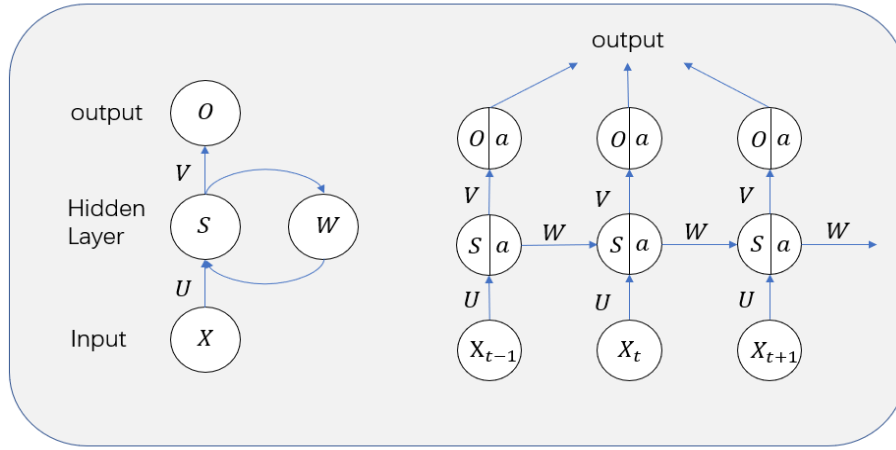


Figure 3.8: A Simple RNN structure

Recurrent Neural Network is one of the most important networks in this project, it provides us with the basic idea to build a better model. Recall from the last section, we can do some classification by using the MLP. In the real world, there are some data with sequence information that triggers our interest. For example, the stock price is affected by time, we can generate a sequence of the model on different days and the previous price will influence the price in the latter days. Or, we can think about the text, there is also some sequential pattern in our human language.

For instance, a simple sentence in the clinical evidence "The patients will take one dose of drag". From this sentence, we can see the first meaningful word is "patients", and the second one is "take". From the first two words, we can think about the possible word for the next one, "drag" is a reason that appears in a clinical document. Therefore, it is meaningful to use some model that can remember the sequence in the text to learn this text better.

From figure 3.8, we can see that there are two kinds of networks. The left one is a simpler model to show how an RNN works. From the input X to the hidden layer S . Then, recurrently calculate by weight matrix W and finally get the output. The right one is a flatter version of the previous one, it shows a sequence for $t-1$ to $t+1$, but the idea is exactly the same as the previous one. For the forward propagation and the backward propagation procedures, we can refer to the steps in MLP. The main idea is the same, except for an important note that we need to consider. Since the latest one is always based on the previous one, we need to notice the sequence effect when doing the calculation.

3.4.3 Gate Recurrent Unit

Gate Recurrent Unit is a kind of Recurrent Neutron Network. We will choose GRU as part of our model in the classification process. The purpose of the GRU is to low down the computation cost but keep a relatively good result. The basic structure of GRU is shown in the left part of Figure 3.9. The h_{t-1} and h_t are the hidden state in $t-1$ and t respectively. X_t and Y_t are the input and output at time t .

If we look inside the GRU model, we can have a better understanding

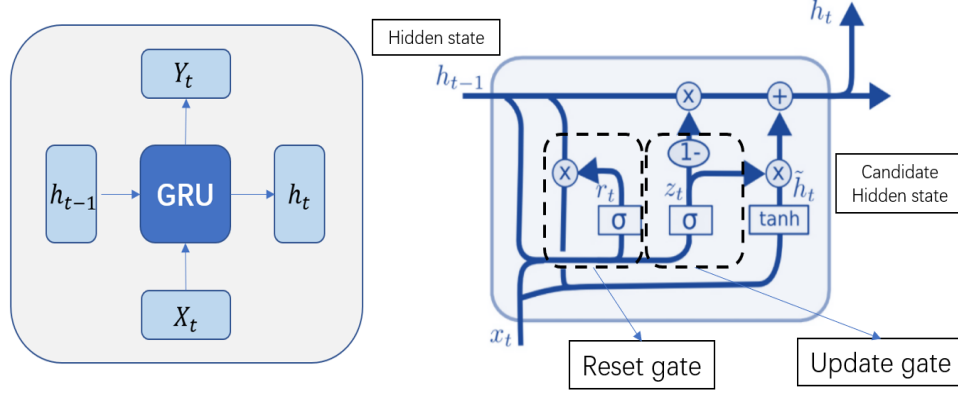


Figure 3.9: The structure of GRU

of how it works. For the right part of the figure, we can see that there are two gates called reset gate and update gate in the structure. We define the r to be the reset gate, where $r = \sigma(W_r \cdot [h_{t-1}, x_t])$, and z to be the update gate, where $z = \sigma(W_z \cdot [h_{t-1}, x_t])$. To reset the h_{t-1} , we need to use the equation $h'_{t-1} = h_{t-1} \odot r$, and then use the activation function $\tanh(W'_h \cdot [r \cdot h_{t-1}, x_t])$ to get h' , where it contains the information from the last hidden state and contain the information from the new input X_t . For the update steps, the updated h_t should equal to $(1 - z) \odot h_{t-1} + z \odot h'$. To understand this equation, we need to know that the range of z is from 0 to 1, which means the updated h_t is actually keep the $1 - z$ proportion from the last hidden state h_{t-1} and add the z proportion from the h' . To explain it further, it is some kind of balance between the "forget" and "remember", where we will remember all the things if z is 1. The strength of this kind of mechanism is that we can use only one update gate z to remember and forget the information at the same time, and the information proportion is under our control.

For the advantages, the proposed GRU will fix the vanishing gradient problem due to its gate structure. In the traditional RNN model, there may exist exploding gradient problem or vanishing gradient problem if there are too many hidden states in the structure. Gradient exploding and vanishing means the gradient will become too large or even equal to zero respectively. In this case, the model will no longer be in a learning status. In the structure of GRU, since the gates structure can help the model to have a memory of the previous hidden state, the sequence dependency problems that appeared in RNN will be fixed.

3.4.4 Bidirectional Layers

The existence of bidirectional layers is to apply the original model twice in different directions and combine them together. Traditionally, an RNN model will be applied to follow the exact sequence of the text, while a bidirectional layer adds a path from the end to the beginning of the text so that we can dive deeper into the text by providing a whole picture for the context. The basic structure is shown in figure 3.10.

3.4.5 Attention context

Attention context is another technique that will be added to our project. In natural language processing, there are many kinds of attention models. In this section, we will give a brief introduction by giving a general structure of attention that is used that combines the GRU model. The ultimate purpose of this mechanism is to help us focus more on valuable information and ignore some useless parts by giving different weights to the data. Generally speaking, there are three steps in total. First, we need




Figure 3.10: RNN with bidirectional layer

to calculate the similarity of query and keys. Then, the similarity output should be normalized by using the softmax function. Finally, calculate the summation of the multiplication of values and the similarity after softmax to get the attention value.

$$Attention(Query, Source) = \sum_{i=1}^L \text{Similarity}(Query, Key_i) \cdot Values_i \quad (3.8)$$

For the advantages, when we add the attention layer on our GRU model, it will give weight to different hidden states, which is very useful for focusing on the more important part of the text and will increase our model performance. On the contrary, if we only use the GRU model to do

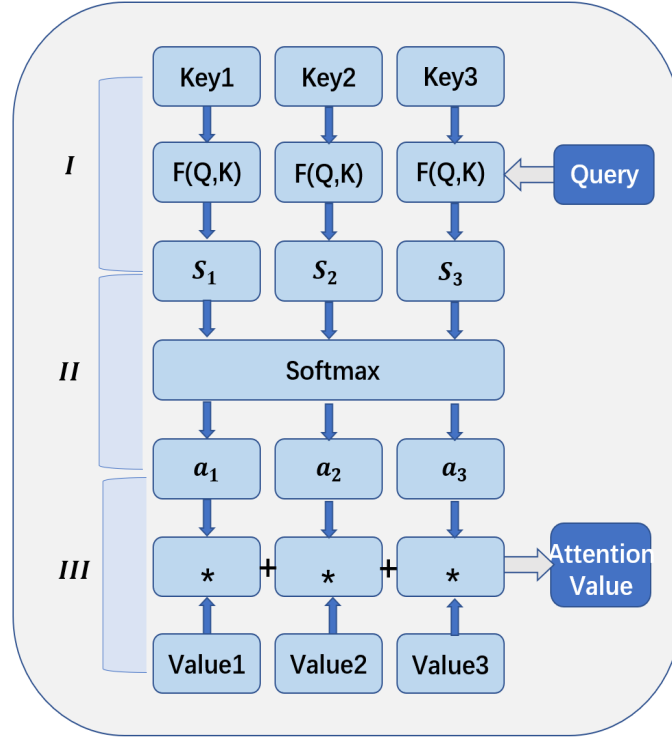


Figure 3.11: The structure of Attention Context

the classification, some of the information will be lost when we consider all the hidden states with equivalent importance.

3.4.6 Final Proposed Model

All the introduced theoretical concepts of the model provide us with great basics to construct a better model. Therefore, a GRU with a Bidirectional layer and attention layer added model was finally proposed. The overview structure is shown in figure 3.12, and we will apply this model to the classification process in the later chapter.

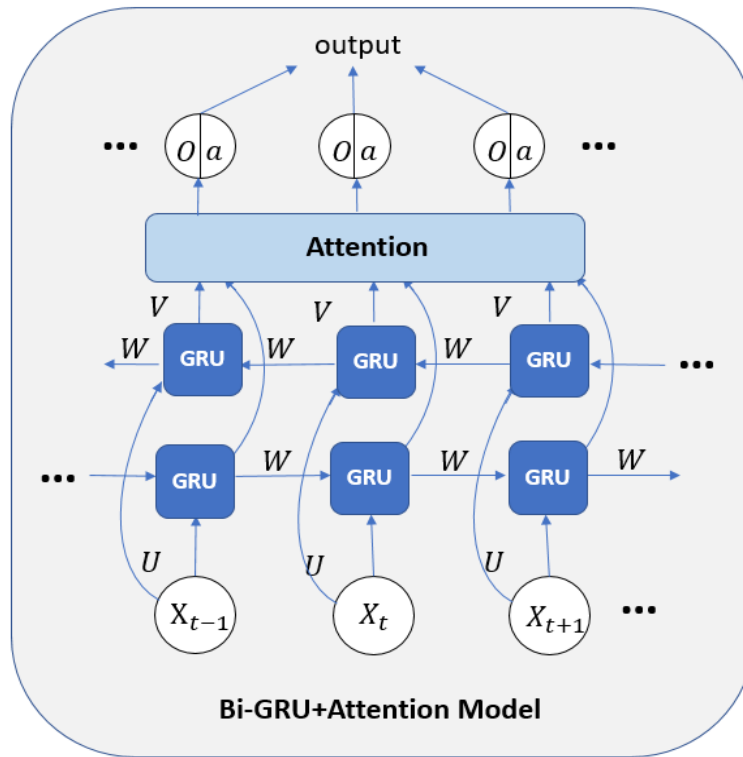


Figure 3.12: Structure of Bi-GRU+Attention Model

Chapter 4

Experiment and Results

4.1 Word2Vec + LightGBM

For our proposed Model 1, first of all, we use Word2Vec Model (we choose CBOW) from Gensim to get the vector representation of the words. Then, by applying the average feature function to correct the input size of the classification model. Then, after setting the parameters, we use LightGBM as the classification model to get the predicted results.

The parameter setting of the word2vec model, average features function, and LightGBM are listed in Table 4.1 and Table 4.2. Also, the training steps and model evaluation are shown in Table 4.3 and Table 4.4. Finally, we got a 60% accuracy for the predicted in the label of training data. After submission on Kaggle, we got a private score: 3.59873, a public score: 1.46629, and the final rank of the private score is 302 out of 1386.

Model Parameters - Word2Vec			
Model	CBOW	Model	Skip-Gram
min_count	1	min_count	1
vector size	200	vector size	200
window	5	window	5
		sg	1
Model Parameters - Average Features			
sentence	train	num_features	200
model	word2vec.wv		

Table 4.1: Parameters setting for Word2Vec and Average Features

Model Parameters - LightGBM			
boosting_type	gbdt	feature_fraction	0.9
objective	multiclass	bagging_fraction	0.8
num_class	9	bagging_freq	5
metric	multi_error	lambda_l1	0.4
num_leaves	500	lambda_l2	0.5
min_data_in_leaf	100	min_gain_to_split	0.2
learning_rate	0.1	verbose	-1

Table 4.2: Parameters setting for LightGBM

Training	
Model1: Train data length: 2324	
Model1: Test data length: 997	
Training until validation scores don't improve for 300 rounds	
[50]	valid_0's multi_error: 0.437312
[100]	valid_0's multi_error: 0.418255
[150]	valid_0's multi_error: 0.417252
[200]	valid_0's multi_error: 0.418255
[250]	valid_0's multi_error: 0.412237
[300]	valid_0's multi_error: 0.417252
[350]	valid_0's multi_error: 0.415246
[400]	valid_0's multi_error: 0.416249
[450]	valid_0's multi_error: 0.406219
[500]	valid_0's multi_error: 0.406219
[550]	valid_0's multi_error: 0.402207
[600]	valid_0's multi_error: 0.408225
[650]	valid_0's multi_error: 0.402207
[700]	valid_0's multi_error: 0.410231
[750]	valid_0's multi_error: 0.408225
Early stopping, best iteration is:	
[491]	valid_0's multi_error: 0.398195
all tasks done. total time used:3.629003 s.	
auc 0.4502928863833586	

Table 4.3: Training for LightGBM

Model Evaluation				
	precision	recall	f1-score	support
0	0.53	0.58	0.55	170
1	0.55	0.40	0.47	146
2	0.44	0.29	0.35	28
3	0.65	0.64	0.64	208
4	0.43	0.36	0.39	72
5	0.82	0.67	0.74	73
6	0.62	0.77	0.69	287
7	0.00	0.00	0.00	3
8	1.00	0.50	0.67	10
accuracy			0.60	997
macro avg	0.56	0.47	0.50	997
weighted avg	0.60	0.60	0.59	997

Table 4.4: Model Evaluation for LightGBM

4.2 Word Embedding + Bi-GRU, Attention

For our proposed Model 2, first of all, we use Tokenizer from Kears to pad the texts to sequence for each observation. And do word embedding to make the input become a three-dimension tenor with matrix representation. Then, training with our proposed bi-directional GRU models and add attention layers to get the predicted output.

The parameter setting of the word embedding and Bidirectional GRU with Attention Model are listed in Table 4.5. Also, the model structure and training steps are shown in Table 4.6 and Table 4.7. Finally, we got an 84.52% accuracy for the predicted in the label of training data. After submission on Kaggle, we get a private score: 2.52648, public score: 1.76996, and the final rank of the private score is 91 out of 1386 when the length of the sequence parameter is 1000. Later, we found that by

Model Parameters - Bi-GRU with Attention Model			
NUM_CLASS	9	Epoch	10
VOCABULARY_SIZE	10000	Batch size	32
SEQUENCE_LENGTH	1000/2000	Val_Epoch	3
Embedding_dim	200	Val_Batch size	32
lstm_out	64	recurrent_dropout	0.2
dense	32	dropout	0.2

Table 4.5: Parameters for Bi-GRU with Attention Model

increasing the sequence length, we will get a little bit better result in Kaggle. When increase the sequence length from 1000 to 2000, we got the private score: 2.36962, public score: 1.88202, and the final private rank is 88 out of 1386.

Bidirectional GRU with Attention Model			
Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None,1000)]	0	[]
input_2 (InputLayer)	[(None,1000)]	0	[]
embedding (Embedding)	(None, 1000,200)	2000000	['input_1 [0][0]']
embedding_1 (Embedding)	(None, 1000,200)	2000000	['input_2 [0][0]']
input_3 (InputLayer)	[(None,1522)]	0	[]
input_4 (InputLayer)	[(None,347)]	0	[]
bidirectional (Bidirectional)	(None, 1000, 128)	102144	['embedding [0][0]']
bidirectional_1 (Bidirectional)	(None, 1000, 128)	102144	['embedding_1 [0][0]']
concatenate (Concatenate)	(None,1869)	0	['input_3[0][0]', 'input_4[0][0]']
attention_with_context (AttentionWithContext)	(None,128)	16640	['bidirectional [0][0]']
attention_with_context_1 (AttentionWithContext)	(None,128)	16640	['bidirectional_1 [0][0]']
dense (Dense)	(None,32)	59840	['concatenate [0][0]']
concatenate_1 (Concatenate)	(None,288)	0	['attention_with_ context[0][0]', 'attention_with_ context_1[0][0]', 'dense[0][0]']
dense_1 (Dense)	(None,9)	2601	['concatenate_1[0][0]']
Total params: 4,300,009			
Trainable params: 4,300,009			
Non-trainable params: 0			

Table 4.6: Bi-GRU with Attention Model

Training for Bi-GRU with Attention Model	
Epoch 1/10	
- loss: 1.6689 - accuracy: 0.3797	
Epoch 00001: val_loss improved from inf to 0.08144, saving model to keras_model	
-loss: 1.6689 - accuracy: 0.3797 - val_loss: 0.0814 - val_accuracy: 0.0625	
Epoch 2/10	
- loss: 1.1092 - accuracy: 0.5929	
Epoch 00002: val_loss improved from 0.08144 to 0.06775	
- loss: 1.1092 - accuracy: 0.5929 - val_loss: 0.0678 - val_accuracy: 0.1637	
Epoch 3/10	
- loss: 0.8239 - accuracy: 0.7025	
Epoch 00003: val_loss improved from 0.06775 to 0.06573	
- loss: 0.8239 - accuracy: 0.7025 - val_loss: 0.0657 - val_accuracy: 0.3670	
Epoch 4/10	
- loss: 0.6516 - accuracy: 0.7712	
Epoch 00004: val_loss improved from 0.06573 to 0.06354	
- loss: 0.6516 - accuracy: 0.7712 - val_loss: 0.0635 - val_accuracy: 0.2945	
Epoch 5/10	
- loss: 0.5466 - accuracy: 0.7983	
Epoch 00005: val_loss did not improve from 0.06354	
- loss: 0.5466 - accuracy: 0.7983 - val_loss: 0.0696 - val_accuracy: 0.2096	
Epoch 6/10	
- loss: 0.4893 - accuracy: 0.8118	
Epoch 00006: val_loss did not improve from 0.06354	
- loss: 0.4893 - accuracy: 0.8118 - val_loss: 0.0704 - val_accuracy: 0.1978	
Epoch 7/10	
- loss: 0.4372 - accuracy: 0.8223	
Epoch 00007: val_loss did not improve from 0.06354	
- loss: 0.4372 - accuracy: 0.8223 - val_loss: 0.0731 - val_accuracy: 0.2971	
Epoch 8/10	
- ETA: 0s - loss: 0.4073 - accuracy: 0.8383	
Epoch 00008: val_loss did not improve from 0.06354	
- loss: 0.4073 - accuracy: 0.8383 - val_loss: 0.0720 - val_accuracy: 0.1971	
Epoch 9/10	
- ETA: 0s - loss: 0.3917 - accuracy: 0.8398	
Epoch 00009: val_loss did not improve from 0.06354	
- loss: 0.3917 - accuracy: 0.8398 - val_loss: 0.0759 - val_accuracy: 0.2041	
Epoch 10/10	
- ETA: 0s - loss: 0.3667 - accuracy: 0.8452	
Epoch 00010: val_loss did not improve from 0.06354	
- loss: 0.3667 - accuracy: 0.8452 - val_loss: 0.0765 - val_accuracy: 0.2101	

Table 4.7: Training for Bi-GRU with Attention Model

4.3 Results

Model1 - Random 5 Results out of 986 Results

ID	class1	class2	class3	class4	class5	class6	class7	class8	class9
559	0.596669	0.072487	0.000433	0.027216	0.001185	0.016108	0.278105	0.004147	0.003651
560	0.053238	0.002687	0.041341	0.003291	0.815927	0.082539	0.000829	5.35E-05	9.45E-05
561	0.430737	0.00123	0.001359	0.016429	0.539503	0.008062	0.00171	0.000235	0.000736
562	0.000517	7.87E-05	3.95E-05	0.994216	0.000372	0.000246	0.004159	0.000169	0.000202
563	0.022975	0.211937	0.000588	0.107404	0.052235	0.001869	0.598732	0.002087	0.002173
...									
986									
ID	class1	class2	class3	class4	class5	class6	class7	class8	class9
559	1	0	0	0	0	0	0	0	0
560	0	0	0	0	1	0	0	0	0
561	0	0	0	0	1	0	0	0	0
562	0	0	0	1	0	0	0	0	0
563	0	0	0	0	0	0	1	0	0
...									
986									

Figure 4.1: Classification Results-1

For the classification Results, we can refer to Figure 4.1, 4.2, and 4.3. For the submission results, we make three submissions on Kaggle. For model 1 with the Lightgbm model, we get a private score: 3.59873, a public score: 1.46629, and the final rank of the private score is 302 out of 1386. For model 2 with Bi-GRU and Attention model, we get a private score: 2.52648, public score: 1.76996, and the final rank of the private score is 91 out of 1386 when the length of the sequence parameter is 1000. We get the best model with Bi-GRU and Attention model when the length of the sequence is 2000, where the private score: 2.36962, public score: 1.88202, and the final private rank is 88 out of 1386. (Note: The public score is for testing purposes in the Kaggle competition. The Private score will be considered as the useful score in the competition.)

Model2-1 Random 5 Results out of 986 Results Sequence length 1000

ID	class1	class2	class3	class4	class5	class6	class7	class8	class9
559	0.170628	0.147612	0.053016	0.166229	0.09173	0.083567	0.243275	0.016068	0.027874
560	0.126328	0.045024	0.035029	0.05354	0.639337	0.06412	0.021932	0.008378	0.006313
561	0.210251	0.124376	0.042793	0.159376	0.110343	0.110612	0.215928	0.008554	0.017766
562	0.03284	0.018627	0.032981	0.816597	0.027288	0.037492	0.020632	0.004874	0.008668
563	0.166839	0.14777	0.053195	0.16634	0.094276	0.08245	0.246316	0.016318	0.026495

...

ID	class1	class2	class3	class4	class5	class6	class7	class8	class9
559	0	0	0	0	0	0	1	0	0
560	0	0	0	0	1	0	0	0	0
561	0	0	0	0	0	0	1	0	0
562	0	0	0	1	0	0	0	0	0
563	0	0	0	0	0	0	1	0	0

...

986

Figure 4.2: Classification Results-2

Model2-2 Random 5 Results out of 986 Results Sequence length 2000

ID	class1	class2	class3	class4	class5	class6	class7	class8	class9
559	0.141351	0.149183	0.06552	0.175357	0.074423	0.110917	0.225321	0.028615	0.029313
560	0.056723	0.01398	0.121212	0.027628	0.316461	0.045233	0.313854	0.039469	0.065441
561	0.137214	0.150981	0.06545	0.176246	0.073253	0.106498	0.235477	0.028375	0.026506
562	0.140977	0.152256	0.064073	0.17613	0.074437	0.107476	0.228968	0.028351	0.027334
563	0.138861	0.149943	0.065433	0.174805	0.074679	0.10773	0.232379	0.028695	0.027474

...

ID	class1	class2	class3	class4	class5	class6	class7	class8	class9
559	0	0	0	0	0	0	1	0	0
560	0	0	0	0	1	0	0	0	0
561	0	0	0	0	0	0	1	0	0
562	0	0	0	0	0	0	1	0	0
563	0	0	0	0	0	0	1	0	0

...

986

Figure 4.3: Classification Results-3

FINAL RESULTS FOR THIS PROJECT			
Submission and Description	Private Score	Public Score	Private Rank
Final_Submission_lightGBM.csv a few seconds ago by KaiyangL Final Submission LightGBM	3.59873	1.46629	302/1386
Submission and Description	Private Score	Public Score	
Final_Submission_BiGRUAttention - 1000.csv just now by KaiyangL Final Submission Bi-GRU+Attention - Seq len 1000	2.52648	1.76996	91/1386
The Best Model			
Submission and Description	Private Score	Public Score	Private Rank
Final_Submission_BiGRUAttention - 2000.csv just now by KaiyangL Final Submission Bi-GRU+Attention - Seq len 2000	2.36962	1.88202	88/1386

Figure 4.4: Final Results

Chapter 5

Conclusions

In conclusion, two classification models are involved in our project. For the LightGBM model, we use the word vector from the Word2Vec model and do the feature average as the input, we successfully predict the class label. The strength of this model is very obvious, only a few seconds are needed when training the LightGBM model and at the same time get a result with relatively high accuracy. For the Bidirectional GRU model with an attention layer, we use the three-dimensional tensor as the input. Looking deeper into the model itself, thanks to the advantages of our model in many ways, we got a high accuracy up to around 85%. The gate structure from GRU will make the balance of the memory and new information mode easier and use fewer computation resources. Also, the bidirectional layer enables us to get the full picture of the sequence from two directions of the texts. Moreover, the adding of attention mechanism greatly improves the performance of our model by paying more attention to the useful information in the texts.

For the limitation, we may find an interesting phenomenon in the result given by Kaggle. The LightGBM model has the highest public score among the three submissions but gets the lowest score in the private score. To explain it further, the test data sets for public and private scores are different. The result in LightGBM shows that it can learn the text in a public dataset very well due to the potential over-fitting problems, and it cannot handle the text in a private dataset since the lack of generalization ability. For the Bi-GRU with attentions model, we make full use of the validation set and set the appropriate drop-out parameter to avoid over-fitting. However, for each run of this model, we will need a few hours to get the results. Therefore, those complex models will be greatly restrained by the computation limitation of the computer.

Nowadays, with the great help of machine learning and deep learning model in the classification of clinical documents, especially in dealing with the text that is related to the gene mutation, lots of human efforts will be avoided. In the future, some high-performance model with a lower computational cost is desirable to develop in dealing with the classification of the clinical text so that patients may receive a personalized treatment that benefits from NLP techniques and classification models.

Appendix A

Python Code

```

1  # **FYP Final Version - LightGBM + (Bi-GRU and Attention)**
2
3  # # Packages Loading
4  import re
5  import time
6  import os
7  import math
8  import pandas as pd
9  import numpy as np
10
11 # Preprocess
12 import spacy
13 from scipy.stats import entropy
14 import string
15 import nltk
16 nltk.download('stopwords')
17 from nltk.tokenize import word_tokenize
18 from nltk.corpus import stopwords
19 from nltk.stem import PorterStemmer
20 from nltk import FreqDist
21 nltk.download('punkt')
22 stop_words = set(stopwords.words('english'))
23 nltk.download('vader_lexicon')
24 from nltk.sentiment.vader import SentimentIntensityAnalyzer
25
26 # Gensim
27 import gensim
28 from gensim.models import LdaModel
29 from gensim import models, corpora, similarities
30 from gensim.models.word2vec import Word2Vec
31
32 # Sklearn
33 from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer, TfidfTransformer
34 from sklearn.cluster import MiniBatchKMeans, KMeans
35 from sklearn.decomposition import PCA
36 from sklearn.metrics import homogeneity_score
37 from sklearn.metrics import silhouette_score
38 from sklearn.model_selection import cross_val_predict
39 from sklearn.model_selection import train_test_split
40 from sklearn.linear_model import LogisticRegression
41 from sklearn.metrics import log_loss, accuracy_score
42 from sklearn.svm import SVC
43 from sklearn.decomposition import TruncatedSVD
44 from sklearn.ensemble import RandomForestClassifier
45 from sklearn import preprocessing
46 from sklearn.preprocessing import LabelEncoder
47 import scikitplot.plotters as skplt
48
49 # Modeling
50 import lightgbm as lgb
51 from datetime import datetime
52
53 # Viz
54 get_ipython().run_line_magic('matplotlib', 'inline')
55 import matplotlib.pyplot as plt
56 import seaborn as sns
57
58 # Sklearn
59 from sklearn.datasets import load_digits
60 from sklearn.model_selection import train_test_split
61 from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix

```

```

62 from sklearn.preprocessing import label_binarize
63 from sklearn.metrics import roc_curve, auc
64 from sklearn.metrics import confusion_matrix
65 from sklearn.metrics import roc_auc_score
66
67 # Kears
68 from keras.preprocessing.sequence import pad_sequences
69 from keras.preprocessing.text import Tokenizer
70 from keras.utils import np_utils
71 from keras.layers.merge import concatenate
72 from keras.utils.np_utils import to_categorical
73 from keras.callbacks import ModelCheckpoint
74 from keras.models import load_model
75 from keras.optimizers import adam_v2
76
77 from keras import backend as K
78 import tensorflow.python.keras.engine
79 from tensorflow.python.keras.layers import Layer, InputSpec
80 from tensorflow.keras.layers import Layer, InputSpec
81 from keras import initializers, regularizers, constraints
82
83 # 搭建模型
84 from keras.models import Sequential, Model
85 from keras.layers import Dense, Embedding, Activation, Input
86 from keras.layers import Convolution1D, Flatten, Dropout, MaxPool1D
87 from keras.layers import BatchNormalization
88 from keras.layers import Convolution1D, Conv1D, MaxPooling1D
89 from keras.layers import Dense, Embedding, Input, Lambda, Reshape
90 from keras.layers import Convolution1D, Flatten, Dropout, MaxPool1D, GlobalAveragePooling1D
91 from keras.layers import LSTM, GRU, TimeDistributed, Bidirectional
92 from keras.layers import Dense, Embedding, LSTM, GRU, Bidirectional, merge, Input, concatenate
93 from keras.layers.merge import Concatenate
94
95
96 # # Data Exploration
97
98 # ## Loading Data
99
100 train_variants_df = pd.read_csv("training_variants.csv")
101 train_text_df = pd.read_csv("training_text.zip", sep="\\", engine="python", names=["ID", "Text"], skiprows=1)
102 test_variants_df = pd.read_csv("stage2_test_variants.csv")
103 test_text_df = pd.read_csv("stage2_test_text.csv", sep="\\", engine="python", names=["ID", "Text"], skiprows=1)
104
105 val_variants_df = pd.read_csv("test_variants.csv")
106 val_text_df = pd.read_csv("test_text.zip", sep="\\", engine="python", names=["ID", "Text"], skiprows=1)
107 val_labels_df = pd.read_csv("stage1_solution_filtered.csv")
108 val_labels_df['Class'] = pd.to_numeric(val_labels_df.drop('ID', axis=1).idxmax(axis=1).str[5:])
109 val_labels_df = val_labels_df[['ID', 'Class']]
110 val_text_df = pd.merge(val_text_df, val_labels_df, how='left', on='ID')
111
112 print("Train Variant".ljust(15), train_variants_df.shape)
113 print("Train Text".ljust(15), train_text_df.shape)
114 print("Test Variant".ljust(15), test_variants_df.shape)
115 print("Test Text".ljust(15), test_text_df.shape)
116 print("Validation Variant".ljust(15), val_variants_df.shape)
117 print("Validation Text".ljust(15), val_text_df.shape)
118
119 train_variants_df.head()
120 test_variants_df.head()
121 val_variants_df.head()

```

```

122 train_variants_df['Class'].value_counts()
123 train_text_df
124 test_text_df
125 val_text_df
126
127 # ## Distribution of genetic mutation classes
128
129 plt.figure(figsize=(14,8))
130 sns.countplot(x="Class", data=train_variants_df, palette="PuBu")
131 plt.ylabel('Frequency', fontsize=22)
132 plt.xlabel('Classes of Genes Variation', fontsize=22)
133 plt.title("Distribution of Genetic Mutation Classes", fontsize=24)
134 plt.show()
135
136 gene_group = train_variants_df.groupby("Gene")['Gene'].count()
137 minimal_occ_genes = gene_group.sort_values(ascending=True)[:10]
138 print("Genes with maximal occurrences\n", gene_group.sort_values(ascending=False)[:10])
139 print("\nGenes with minimal occurrences\n", minimal_occ_genes)
140
141
142 # ## Distribution of Gene in Different Classes
143
144 fig, axs = plt.subplots(ncols=3, nrows=3, figsize=(20,15))
145
146 for i in range(3):
147     for j in range(3):
148         gene_count_grp = train_variants_df[train_variants_df["Class"]==((i*3+j)+1)]
149         .groupby('Gene')['ID'].count().reset_index()
150         sorted_gene_group = gene_count_grp.sort_values('ID', ascending=False)
151         sorted_gene_group_top_7 = sorted_gene_group[:7]
152         sns.barplot(x="Gene", y="ID", data=sorted_gene_group_top_7, ax=axs[i][j], palette="PuBu")
153
154
155 # Some points we can conclude from these graphs:
156 #
157 # BRCA1 is highly dominating Class 5\
158 # SF3B1 is highly dominating Class 9\
159 # BRCA1 and BRCA2 are dominating Class 6
160
161 # # Data Preprocessing
162
163 # ## Text Preprocessing
164
165 # **Steps**\
166 # **1. Tokenization**\
167 # **2. Removal of punctuations**\
168 # **3. Lemmatization**\
169 # **4. Removal of stop words**\
170 # **5. Lower casting**\
171 # **6. Special consideration for clinical text**
172
173 # Remove punctuation
174 train_text_df['Text_processed'] = train_text_df['Text'].map(lambda x: re.sub('[\',.!?*]', '', str(x)))
175 # Convert the titles to lowercase
176 train_text_df['Text_processed'] = train_text_df['Text_processed'].replace(r'\n', ' ', regex=True)
177 train_text_df['Text_processed'] = train_text_df['Text_processed'].map(lambda x: x.lower())
178 # Print out the first rows of papers
179 train_text_df['Text_processed'] = train_text_df['Text_processed'].apply(lambda x: x.strip())
180 train_text_df
181

```

```

182 # Remove punctuation
183 test_text_df['Text_processed'] = test_text_df['Text'].map(lambda x: re.sub('[\',.!?*]', '', str(x)))
184 # Convert the titles to lowercase
185 test_text_df['Text_processed'] = test_text_df['Text_processed'].replace(r'\n', ' ', regex=True)
186 test_text_df['Text_processed'] = test_text_df['Text_processed'].map(lambda x: x.lower())
187 # Print out the first rows of papers
188 test_text_df['Text_processed'] = test_text_df['Text_processed'].apply(lambda x: x.strip())
189 test_text_df
190
191 # Remove punctuation
192 val_text_df['Text_processed'] = val_text_df['Text'].map(lambda x: re.sub('[\',.!?*]', '', str(x)))
193 # Convert the titles to lowercase
194 val_text_df['Text_processed'] = val_text_df['Text_processed'].replace(r'\n', ' ', regex=True)
195 val_text_df['Text_processed'] = val_text_df['Text_processed'].map(lambda x: x.lower())
196 # Print out the first rows of papers
197 val_text_df['Text_processed'] = val_text_df['Text_processed'].apply(lambda x: x.strip())
198 val_text_df
199
200 train_full = train_variants_df.merge(train_text_df, how="inner", left_on="ID", right_on="ID")
201 train_full = train_full.drop("Text", axis=1)
202 train_full.head()
203
204 test_full = test_variants_df.merge(test_text_df, how="inner", left_on="ID", right_on="ID")
205 test_full = test_full.drop("Text", axis=1)
206 test_full.head()
207
208 val_full = val_variants_df.merge(val_text_df, how="inner", left_on="ID", right_on="ID")
209 val_full = val_full.drop("Text", axis=1)
210 val_full.head()
211
212 import nltk
213 nltk.download('words')
214 words = set(nltk.corpus.words.words())
215 train_full['Text_processed'] = train_full['Text_processed'].apply(lambda x: " "
216 | .join(w for w in nltk.wordpunct_tokenize(x) if w.lower() in words ))
217 test_full['Text_processed'] = test_full['Text_processed'].apply(lambda x: " "
218 | .join(w for w in nltk.wordpunct_tokenize(x) if w.lower() in words ))
219 val_full['Text_processed'] = val_full['Text_processed'].apply(lambda x: " "
220 | .join(w for w in nltk.wordpunct_tokenize(x) if w.lower() in words ))
221
222 ### Word Statistics
223
224 train_text_df.loc[:, 'Text_count'] = train_text_df["Text_processed"].apply(lambda x: len(x.split()))
225 train_text_df.head()
226 test_text_df.loc[:, 'Text_count'] = test_text_df["Text_processed"].apply(lambda x: len(x.split()))
227 test_text_df.head()
228 val_text_df.loc[:, 'Text_count'] = val_text_df["Text_processed"].apply(lambda x: len(x.split()))
229 val_text_df.head()
230 train_full = train_variants_df.merge(train_text_df, how="inner", left_on="ID", right_on="ID")
231 train_full.head()
232 test_full = test_variants_df.merge(test_text_df, how="inner", left_on="ID", right_on="ID")
233 test_full.head()
234 val_full = val_variants_df.merge(val_text_df, how="inner", left_on="ID", right_on="ID")
235 val_full.head()
236
237 train_full = train_full.drop("Text", axis=1)
238 test_full = test_full.drop("Text", axis=1)
239 val_full = val_full.drop("Text", axis=1)
240
241 print(sum(train_full["Text_count"]))

```

```

242 print(sum(test_full["Text_count"]))
243 print(sum(val_full["Text_count"]))
244
245
246 from wordcloud import WordCloud
247 # Join the different processed titles together.
248 long_string = ','.join(list(train_text_df['Text_processed'].values))
249 # Create a WordCloud object
250 wordcloud = WordCloud(background_color="black", max_words=5000, contour_width=5, contour_color='steelblue')
251 # Generate a word cloud
252 wordcloud.generate(long_string)
253 # Visualize the word cloud
254 wordcloud.to_image()
255
256 count_grp = train_full.groupby('Class')['Text_count']
257 count_grp.describe()
258
259 plt.figure(figsize=(12,8))
260 gene_count_grp = train_full.groupby('Gene')['Text_count'].sum().reset_index()
261 sns.violinplot(x="Class", y="Text_count", data=train_full, inner=None,palette="Spectral")
262 sns.swarmplot(x="Class", y="Text_count", data=train_full, color="w", alpha=.5);
263 plt.ylabel('Text Count', fontsize=14)
264 plt.xlabel('Class', fontsize=14)
265 plt.title("Text length distribution", fontsize=18)
266 plt.show()
267
268
269 fig, axs = plt.subplots(ncols=3, nrows=3, figsize=(20,16))
270
271 for i in range(3):
272     for j in range(3):
273         gene_count_grp = train_full[train_full["Class"]==((i*3+j)+1)].groupby('Gene')['Text_count'].mean().reset_index()
274         sorted_gene_group = gene_count_grp.sort_values('Text_count', ascending=False)
275         sorted_gene_group_top_7 = sorted_gene_group[:7]
276         sns.barplot(x="Gene", y="Text_count", data=sorted_gene_group_top_7, ax=axs[i][j],palette="Spectral")
277
278
279 # # Vector Representation
280
281 train_full["Text_processed"]
282 test_full["Text_processed"]
283 val_full["Text_processed"]
284
285 # ## Word2Vec
286
287 # ### CBOW and Skip Gram
288
289
290 # Python program to generate word vectors using Word2Vec
291
292 # importing all necessary modules
293 from nltk.tokenize import sent_tokenize, word_tokenize
294 import warnings
295
296 warnings.filterwarnings(action = 'ignore')
297
298 import gensim
299 from gensim.models import Word2Vec
300
301 Word2Vec_dim = 200
302 train = train_full["Text_processed"].to_string()

```



```

303 test = test_full["Text_processed"].to_string()
304 val = val_full["Text_processed"].to_string()
305
306 # Replaces escape character with space
307 f = train.replace("\n", " ")
308
309 data = []
310
311 # iterate through each sentence in the file
312 for i in sent_tokenize(f):
313     temp = []
314     # tokenize the sentence into words
315     for j in word_tokenize(i):
316         temp.append(j.lower())
317
318     data.append(temp)
319
320
321 # Replaces escape character with space
322 f2 = test.replace("\n", " ")
323
324 data2 = []
325
326 # iterate through each sentence in the file
327 for i in sent_tokenize(f2):
328     temp2 = []
329     # tokenize the sentence into words
330     for j in word_tokenize(i):
331         temp2.append(j.lower())
332
333     data2.append(temp2)
334
335
336 # Replaces escape character with space
337 f3 = val.replace("\n", " ")
338
339 data3 = []
340
341 # iterate through each sentence in the file
342 for i in sent_tokenize(f3):
343     temp3 = []
344     # tokenize the sentence into words
345     for j in word_tokenize(i):
346         temp3.append(j.lower())
347
348     data3.append(temp3)
349
350 # Create CBOW model
351 model1 = gensim.models.Word2Vec(data, min_count = 1, vector_size = Word2Vec_dim, window = 5)
352 model2 = gensim.models.Word2Vec(data2, min_count = 1, vector_size = Word2Vec_dim, window = 5)
353 model3 = gensim.models.Word2Vec(data3, min_count = 1, vector_size = Word2Vec_dim, window = 5)
354
355
356 #Create Skip Gram model
357 model4 = gensim.models.Word2Vec(data, min_count = 1, vector_size = Word2Vec_dim, window = 5, sg = 1)
358 model5 = gensim.models.Word2Vec(data2, min_count = 1, vector_size = Word2Vec_dim, window = 5, sg = 1)
359 # model6 = gensim.models.Word2Vec(data2, min_count = 1, vector_size = Word2Vec_dim, window = 5, sg = 1)
360
361 # model1.save("word2vec.model1_CBOw")
362 # model2.save("word2vec.model2_CBOw")

```

```

363 # model3.save("word2vec.model3_CBoW")
364 # model4.save("word2vec.model4_SkipGram")
365 # model5.save("word2vec.model5_SkipGram")
366 # model6.save("word2vec.model6_SkipGram")
367
368
369 # ### Average feature vector
370
371
372 def avg_feature_vector(sentence, model, num_features):
373     words = sentence.replace('\n', " ").replace(',', ' ').replace('.', " ").split()
374     feature_vec = np.zeros((num_features,), dtype="float32")
375     i=0
376     for word in words:
377         try:
378             feature_vec = np.add(feature_vec, model[word])
379         except KeyError as error:
380             feature_vec
381             i = i + 1
382     if Len(words) > 0:
383         feature_vec = np.divide(feature_vec, Len(words)- i)
384     return feature_vec
385
386
387 train_word2vec1 = np.zeros((Len(train_full), Word2Vec_dim), dtype="float32")
388 test_word2vec1 = np.zeros((Len(test_full), Word2Vec_dim), dtype="float32")
389 val_word2vec1 = np.zeros((Len(val_full), Word2Vec_dim), dtype="float32")
390
391 for i in range(Len(train_full)):
392     train_word2vec1[i] = avg_feature_vector(train_full["Text_processed"][i], model1.wv, Word2Vec_dim)
393
394 for i in range(Len(test_full)):
395     test_word2vec1[i] = avg_feature_vector(test_full["Text_processed"][i], model1.wv, Word2Vec_dim)
396
397 for i in range(Len(val_full)):
398     val_word2vec1[i] = avg_feature_vector(val_full["Text_processed"][i], model1.wv, Word2Vec_dim)
399
400
401 train_word2vec2 = np.zeros((Len(train_full), Word2Vec_dim), dtype="float32")
402 test_word2vec2 = np.zeros((Len(test_full), Word2Vec_dim), dtype="float32")
403 val_word2vec2 = np.zeros((Len(val_full), Word2Vec_dim), dtype="float32")
404
405 for i in range(Len(train_full)):
406     train_word2vec2[i] = avg_feature_vector(train_full["Text_processed"][i], model4.wv, Word2Vec_dim)
407
408 for i in range(Len(test_full)):
409     test_word2vec2[i] = avg_feature_vector(test_full["Text_processed"][i], model4.wv, Word2Vec_dim)
410
411 for i in range(Len(val_full)):
412     val_word2vec2[i] = avg_feature_vector(val_full["Text_processed"][i], model4.wv, Word2Vec_dim)
413
414
415 train_word2vec1 = pd.DataFrame(train_word2vec1)
416 train_word2vec1
417 test_word2vec1 = pd.DataFrame(test_word2vec1)
418 test_word2vec1
419 val_word2vec1 = pd.DataFrame(val_word2vec1)
420 val_word2vec1
421 train_word2vec2 = pd.DataFrame(train_word2vec2)
422 train_word2vec2

```

```

423 test_word2vec2 = pd.DataFrame(test_word2vec2)
424 test_word2vec2
425 val_word2vec2 = pd.DataFrame(val_word2vec2)
426 val_word2vec2
427
428
429 # # Modeling
430
431 # ## LightGBM
432
433 lbl = preprocessing.LabelEncoder()
434 data1 = train_word2vec1
435 target = lbl.fit_transform(train_full["Class"].astype(str)) #将提示的包含错误数据类型这一列进行转换
436 X_train1, X_test1, y_train1, y_test1 = train_test_split(data1, target, test_size=0.3, random_state = 42)
437 print("Model1: Train data length:", len(X_train1))
438 print("Model1: Test data length:", len(X_test1))
439 btime = datetime.now()
440 lgb_train1 = lgb.Dataset(X_train1, y_train1)
441 lgb_eval1 = lgb.Dataset(X_test1, y_test1, reference=lgb_train1)
442
443 params = {
444     'task': 'train',
445     'boosting_type': 'gbdt',
446     'objective': 'multiclass',
447     'num_class': 9,
448     'metric': 'multi_error',
449     'num_leaves': 500,
450     'min_data_in_leaf': 100,
451     'learning_rate': 0.1,
452     'feature_fraction': 0.9,
453     'bagging_fraction': 0.8,
454     'bagging_freq': 5,
455     'lambda_l1': 0.4,
456     'lambda_l2': 0.5,
457     'min_gain_to_split': 0.2,
458     'verbose': -1,
459 }
460
461 gbm1 = lgb.train(params, lgb_train1, num_boost_round=1000, valid_sets=lgb_eval1,
462                 verbose_eval = 50, early_stopping_rounds=300)
463 print('all tasks done. total time used:%s s.\n\n'%((datetime.now() - btime).total_seconds()))
464 gbm1.save_model('model1_CBOW.txt')
465 gbm1 = lgb.Booster(model_file='model1_CBOW.txt')
466 y_pred_pa1 = gbm1.predict(X_test1)
467 y_test_oh1 = label_binarize(y_test1, classes= [1,2,3,4,5,6,7,8,9])
468 #y_pred_lightGBM1 = [list(x).index(max(x)) for x in y_prob1]
469 print('auc: ', roc_auc_score(y_test_oh1, y_pred_pa1, average='micro'))
470
471 y_pred1 = y_pred_pa1.argmax(axis=1)
472 confusion_matrix(y_test1, y_pred1)
473
474 precision_score(y_test1, y_pred1, average='micro')
475 recall_score(y_test1, y_pred1, average='micro')
476 f1_score(y_test1, y_pred1, average='micro')
477
478 print(classification_report(y_test1, y_pred1))
479
480 pred_lightgbm_cbow = gbm1.predict(test_word2vec1)
481 #pred_lightgbm_cbow = [list(x).index(max(x)) for x in pred_lightgbm_cbow]
482 #pred_lightgbm_cbow = pd.get_dummies(np.array(pred_lightgbm_cbow) + 1)

```

```

483 pred_lightgbm_cbow
484
485 # pred_lightgbm_skg = gbm2.predict(test_word2vec2)
486 # # pred_lightgbm_skg = [list(x).index(max(x)) for x in pred_lightgbm_skg]
487 # # pred_lightgbm_skg = pd.get_dummies(np.array(pred_lightgbm_skg) + 1)
488 # pred_lightgbm_skg
489
490
491 # ## RNN
492
493
494
495 NUM_CLASS=9
496 VOCABULARY_SIZE = 10000
497 SEQUENCE_LENGTH= 2000
498 tokenizer = Tokenizer(filters='!"#%&()*+,-./:;<=>@[\\]^_`{|}~\t\n', lower=True, split=" ")
499 tokenizer.fit_on_texts(train_full["Text_processed"])
500 vocab = tokenizer.word_index
501
502 # Training
503 training = train_full.sample(frac=1) # shuffle data first
504 training_input = tokenizer.texts_to_sequences(training['Text_processed'].astype(str))
505 training_input_r = [list(reversed(x)) for x in training_input]
506 training_input_begin = pad_sequences(training_input, maxlen=SEQUENCE_LENGTH)
507 training_input_end = pad_sequences(training_input_r, maxlen=SEQUENCE_LENGTH)
508 training_output = pd.get_dummies(training['Class']).values
509
510 # Testing
511 testing_input = tokenizer.texts_to_sequences(test_full['Text_processed'].astype(str))
512 testing_input_r = [list(reversed(x)) for x in testing_input]
513 testing_input_begin = pad_sequences(testing_input, maxlen=SEQUENCE_LENGTH)
514 testing_input_end = pad_sequences(testing_input_r, maxlen=SEQUENCE_LENGTH)
515
516 # Validation
517 val_input = tokenizer.texts_to_sequences(val_full['Text_processed'].astype(str))
518 val_input_r = [list(reversed(x)) for x in val_input]
519 val_input_begin = pad_sequences(val_input, maxlen=SEQUENCE_LENGTH)
520 val_input_end = pad_sequences(val_input_r, maxlen=SEQUENCE_LENGTH)
521 val_output = pd.get_dummies(val_full['Class']).values
522
523 print("Training set shape:", training_input_begin.shape, training_input_end.shape, training_output.shape)
524 print("Testing set shape:", testing_input_begin.shape, testing_input_end.shape)
525 print("Validation set shape:", val_input_begin.shape, val_input_end.shape, val_output.shape)
526
527 # Add gene and variation to predictor
528 gene_label = LabelEncoder()
529 ALL_Genes = np.concatenate([train_full['Gene'], val_full['Gene'], test_full['Gene']])
530 ALL_Variations = np.concatenate([train_full['Variation'], val_full['Variation'], test_full['Variation']])
531 ALL_Variations = np.asarray([v[0]+v[-1] for v in ALL_Variations])
532 print ("The number of unique genes: ", len(np.unique(ALL_Genes)))
533 print ("The number of unique variations:", len(np.unique(ALL_Variations)))
534
535
536 len_train = len(training_input)
537 len_validation = len(val_input)
538 len_test = len(testing_input)
539 print("The length of training input:", len_train)
540 print("The length of testing input:", len_test)
541 print("The length of validation input:", len_validation)
542
543
544 gene_encoded = pd.get_dummies(ALL_Genes).values

```

[illegible]

```

605         constraint=self.u_constraint)
606
607         super(AttentionWithContext, self).build(input_shape)
608
609     def compute_mask(self, input, input_mask=None):
610         # do not pass the mask to the next layers
611         return None
612
613     def call(self, x, mask=None):
614         uit = dot_product(x, self.W)
615
616         if self.bias:
617             uit += self.b
618
619         uit = K.tanh(uit)
620         ait = dot_product(uit, self.u)
621
622         a = K.exp(ait)
623
624         # apply mask after the exp. will be re-normalized next
625         if mask is not None:
626             # Cast the mask to floatX to avoid float64 upcasting in theano
627             a *= K.cast(mask, K.floatx())
628
629         # in some cases especially in the early stages of training the sum may be almost zero
630         # and this results in NaN's. A workaround is to add a very small positive number  $\epsilon$  to the sum.
631         # a /= K.cast(K.sum(a, axis=1, keepdims=True), K.floatx())
632         a /= K.cast(K.sum(a, axis=1, keepdims=True) + K.epsilon(), K.floatx())
633
634         a = K.expand_dims(a)
635         weighted_input = x * a
636         return K.sum(weighted_input, axis=1)
637
638     def compute_output_shape(self, input_shape):
639         return input_shape[0], input_shape[-1]
640
641 #The number of unique genes: 1522
642 #The number of unique variations: 347
643
644
645 # ### GRU + Bidirectional
646
647 Embedding_dim = 200
648 lstm_out = 64
649
650 # Model saving callback
651 ckpt_callback = ModelCheckpoint('keras_model', monitor='val_loss', verbose=1, save_best_only=True, mode='auto')
652
653 input_sequence_begin = Input(shape=(training_input_begin.shape[1],))
654 input_sequence_end = Input(shape=(training_input_end.shape[1],))
655 input_gene = Input(shape=(training_input_gene.shape[1],))
656 input_variant = Input(shape=(training_input_variation.shape[1],))
657
658 merged = concatenate([input_gene, input_variant])
659 dense = Dense(32, activation='sigmoid')(merged)
660
661 embeds_begin = Embedding(VOCABULARY_SIZE, Embedding_dim, input_length = SEQUENCE_LENGTH)(input_sequence_begin)
662 embeds_out_begin = Bidirectional(GRU(lstm_out, recurrent_dropout=0.2, dropout=0.2, return_sequences=True))(embeds_begin)
663 attention_begin = AttentionWithContext()(embeds_out_begin)
664

```

```

665 embeds_end = Embedding(VOCABULARY_SIZE, Embedding_dim, input_length = SEQUENCE_LENGTH)(input_sequence_end)
666 embeds_out_end = Bidirectional(GRU(lstm_out, recurrent_dropout=0.2, dropout=0.2, return_sequences=True))(embeds_end)
667 attention_end = AttentionWithContext()(embeds_out_end)
668
669 merged2 = concatenate([attention_begin, attention_end, dense])
670 dense2 = Dense(9, activation='softmax')(merged2)
671
672 model_RNN = Model(inputs=[input_sequence_begin, input_sequence_end, input_gene, input_variant], outputs=dense2)
673 model_RNN.compile(loss = 'categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
674 print(model_RNN.summary())
675
676
677 model_RNN.fit([training_input_begin, training_input_end, training_input_gene, training_input_variation],
678             training_output, epochs=10, batch_size=32,
679             validation_data=([val_input_begin, val_input_end, val_input_gene, val_input_variation], val_output),
680             callbacks=[ckpt_callback])
681
682
683 probas = model_RNN.predict([val_input_begin, val_input_end, val_input_gene, val_input_variation])
684 pred_indices = np.argmax(probas, axis=1)
685 classes = np.array(range(1, 10))
686 preds = classes[pred_indices]
687 print('Log loss: {}'.format(log_loss(classes[np.argmax(val_output, axis=1)], probas)))
688 print('Accuracy: {}'.format(accuracy_score(classes[np.argmax(val_output, axis=1)], preds)))
689 skplt.plot_confusion_matrix(classes[np.argmax(val_output, axis=1)], preds)
690
691 model_RNN.fit([
692     np.concatenate([training_input_begin, val_input_begin]),
693     np.concatenate([training_input_end, val_input_end]),
694     np.concatenate([training_input_gene, val_input_gene]),
695     np.concatenate([training_input_variation, val_input_variation]),
696     np.concatenate([training_output, val_output]),
697     epochs=3, batch_size=32, callbacks=[ckpt_callback])
698
699
700 probas = model_RNN.predict([testing_input_begin, testing_input_end, testing_input_gene, testing_input_variation])
701
702 # Submission
703 pred_lightgbm_cbow = pd.DataFrame(pred_lightgbm_cbow)
704 pred_lightgbm_cbow
705 NEWSsubmission_df1 = pred_lightgbm_cbow
706 NEWSsubmission_df1.to_csv("NEWSsubmission_lightGBM.csv", index=False)
707 submission_df = pd.DataFrame(probas, columns=['class'+str(c+1) for c in range(9)])
708 submission_df['ID'] = test_full['ID']
709 submission_df
710 submission_df.to_csv("NEWSsubmission_rnn2000.csv", index=False)

```

Bibliography

- [1] J. Ferlay, M. Ervik, F. Lam, M. Colombet, L. Mery, and M. Piñeros, “Global Cancer Observatory: Cancer Today,” International Agency for Research on Cancer, Tech. Rep., 2021. [Online]. Available: <https://gco.iarc.fr/today>
- [2] J. J. Salk, E. J. Fox, and L. A. Loeb, “Mutational Heterogeneity in Human Cancers: Origin and Consequences,” *Annual Review of Pathology: Mechanisms of Disease*, vol. 5, no. 1, pp. 51–75, 2010. [Online]. Available: <http://www.annualreviews.org/doi/10.1146/annurev-pathol-121808-102113>
- [3] J. Xu, P. Yang, S. Xue, B. Sharma, M. Sanchez-Martin, F. Wang, K. A. Beaty, E. Dehan, and B. Parikh, “Translating cancer genomics into precision medicine with artificial intelligence: applications, challenges and future perspectives,” *Human Genetics*, vol. 138, no. 2, pp. 109–124, 2019. [Online]. Available: <http://link.springer.com/10.1007/s00439-019-01970-5>
- [4] R. Leaman, R. Khare, and Z. Lu, “Challenges in clinical natural language processing for automated disorder normalization,” *Journal of*

- Biomedical Informatics*, vol. 57, pp. 28–37, 2015. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S1532046415001501>
- [5] M. Gupta, H. Wu, S. Arora, A. Gupta, G. Chaudhary, and Q. Hua, “Gene Mutation Classification through Text Evidence Facilitating Cancer Tumour Detection,” *Journal of Healthcare Engineering*, vol. 2021, pp. 1–16, 2021. [Online]. Available: <https://www.hindawi.com/journals/jhe/2021/8689873/>
- [6] M. D. Yandell and W. H. Majoros, “Genomics and natural language processing,” *Nature Reviews Genetics*, vol. 3, no. 8, pp. 601–610, 2002. [Online]. Available: <http://www.nature.com/articles/nrg861>
- [7] S. Velupillai, H. Suominen, M. Liakata, A. Roberts, A. D. Shah, K. Morley, D. Osborn, J. Hayes, R. Stewart, J. Downs, W. Chapman, and R. Dutta, “Using clinical Natural Language Processing for health outcomes research: Overview and actionable suggestions for future advances,” *Journal of Biomedical Informatics*, vol. 88, pp. 11–19, 2018. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S1532046418302016>
- [8] H. Farooq, N. Rehmat, S. Kumar, and H. Naveed, “Genetic Mutation Classification using Machine Learning,” *Biophysical Journal*, vol. 116, no. 3, p. 292a, 2019. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S0006349518328418>
- [9] M. Singh, A. Murthy, and S. Singh, “Prioritization of Free-Text Clinical Documents: A Novel Use of a Bayesian Classifier,” *JMIR Medical Informatics*, vol. 3, no. 2, p. e17, 2015. [Online]. Available: <http://medinform.jmir.org/2015/2/e17/>

- [10] A. Wright, A. B. McCoy, S. Henkin, A. Kale, and D. F. Sittig, “Use of a support vector machine for categorizing free-text notes: assessment of accuracy across two institutions,” *Journal of the American Medical Informatics Association*, vol. 20, no. 5, pp. 887–890, 2013. [Online]. Available: <https://academic.oup.com/jamia/article-lookup/doi/10.1136/amiajnl-2012-001576>
- [11] T. K. Ho, “random decision forests,” *Proceedings of the Third International Conference on Document Analysis and Recognition (Volume 1)*, vol. Volume 1, pp. 278–282, 1995. [Online]. Available: <https://doi.ieeecomputersociety.org/10.1109/ICDAR.1995.598994>
- [12] J. H. Friedman, “Greedy function approximation: A gradient boosting machine.” *The Annals of Statistics*, vol. 29, no. 5, Oct. 2001. [Online]. Available: <https://projecteuclid.org/journals/annals-of-statistics/volume-29/issue-5/Greedy-function-approximation-A-gradient-boosting-machine/10.1214/aos/1013203451.full>
- [13] T. Chen and C. Guestrin, “XGBoost: A Scalable Tree Boosting System,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. San Francisco California USA: ACM, Aug. 2016, pp. 785–794. [Online]. Available: <https://dl.acm.org/doi/10.1145/2939672.2939785>
- [14] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “LightGBM: A Highly Efficient Gradient Boosting Decision Tree,” p. 9.
- [15] X. Qin, J. Liu, Y. Wang, Y. Liu, K. Deng, Y. Ma, K. Zou, L. Li, and X. Sun, “Natural language processing

was effective in assisting rapid title and abstract screening when updating systematic reviews,” *Journal of Clinical Epidemiology*, vol. 133, pp. 121–129, 2021. [Online]. Available: <https://linkinghub.elsevier.com/retrieve/pii/S0895435621000147>

- [16] I. Spasic and G. Nenadic, “Clinical Text Data in Machine Learning: Systematic Review,” *JMIR MEDICAL INFORMATICS*, p. 19.
- [17] S. Hochreiter and J. Schmidhuber, “Long Short-Term Memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997. [Online]. Available: <http://dx.doi.org/10.1162/neco.1997.9.8.1735>
- [18] V. Mnih, N. Heess, and A. Graves, “Recurrent Models of Visual Attention,” p. 9.
- [19] Y. Bao, Z. Deng, Y. Wang, H. Kim, V. D. Armengol, F. Acevedo, N. Ouardaoui, C. Wang, G. Parmigiani, R. Barzilay, D. Braun, and K. S. Hughes, “Using Machine Learning and Natural Language Processing to Review and Classify the Medical Literature on Cancer Susceptibility Genes,” *JCO Clinical Cancer Informatics*, no. 3, pp. 1–9, 2019. [Online]. Available: <https://ascopubs.org/doi/10.1200/CCI.19.00042>
- [20] Y. Wang, S. Sohn, S. Liu, F. Shen, L. Wang, E. J. Atkinson, S. Amin, and H. Liu, “A clinical text classification paradigm using weak supervision and deep representation,” *BMC Medical Informatics and Decision Making*, vol. 19, no. 1, p. 1, 2019. [Online]. Available: <https://bmcmmedinformdecismak.biomedcentral.com/articles/10.1186/s12911-018-0723-6>

- [21] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” *arXiv:1810.04805 [cs]*, 2019. [Online]. Available: <http://arxiv.org/abs/1810.04805>
- [22] Y. Su, H. Xiang, H. Xie, Y. Yu, S. Dong, Z. Yang, and N. Zhao, “Application of BERT to Enable Gene Classification Based on Clinical Evidence,” *BioMed Research International*, vol. 2020, pp. 1–13, 2020. [Online]. Available: <https://www.hindawi.com/journals/bmri/2020/5491963/>
- [23] S. Wu, K. Roberts, S. Datta, J. Du, Z. Ji, Y. Si, S. Soni, Q. Wang, Q. Wei, Y. Xiang, B. Zhao, and H. Xu, “Deep learning in clinical natural language processing: a methodical review,” *Journal of the American Medical Informatics Association*, vol. 27, no. 3, pp. 457–470, 2020. [Online]. Available: <https://academic.oup.com/jamia/article/27/3/457/5651084>